

# Advanced Data Manipulation: Exercises

## Exercise set 1

1. Display the data where the gene ontology biological process (the `bp` variable) is “leucine biosynthesis” (case-sensitive) *and* the limiting nutrient was Leucine. (Answer should return a 24-by-7 data frame – 4 genes  $\times$  6 growth rates).
2. Gene/rate combinations had high expression (in the top 1% of expressed genes)? *Hint:* see `?quantile` and try `quantile(ydat$expression, probs=.99)` to see the expression value which is higher than 99% of all the data, then `filter()` based on that. Try wrapping your answer with a `View()` function so you can see the whole thing. What does it look like those genes are doing? Answer should return a 1971-by-7 data frame.

## Exercise set 2

1. First, re-run the command you used above to filter the data for genes involved in the “leucine biosynthesis” biological process *and* where the limiting nutrient is Leucine.
2. Wrap this entire filtered result with a call to `arrange()` where you’ll arrange the result of #1 by the gene symbol.
3. Wrap this entire result in a `View()` statement so you can see the entire result.

## Exercise set 3

Here’s a warm-up round. Try the following.

1. Show the limiting nutrient and expression values for the gene ADH2 when the growth rate is restricted to 0.05. *Hint:* 2 pipes: `filter` and `select`.

```
## # A tibble: 6 × 2
##   nutrient expression
##   <chr>      <dbl>
## 1 Glucose      6.28
## 2 Ammonia      0.55
## 3 Phosphate   -4.60
## 4 Sulfate     -1.18
## 5 Leucine      4.15
## 6 Uracil       0.63
```

2. What are the four most highly expressed genes when the growth rate is restricted to 0.05 by restricting glucose? Show only the symbol, expression value, and GO terms. *Hint:* 4 pipes: `filter`, `arrange`, `head`, and `select`.

```
## # A tibble: 4 × 4
##   symbol expression      bp                                mf
##   <chr>      <dbl>      <chr>                                <chr>
## 1 ADH2      6.28      fermentation* alcohol dehydrogenase activity
## 2 HSP26     5.86      response to stress* unfolded protein binding
## 3 MLS1      5.64      glyoxylate cycle malate synthase activity
## 4 HXT5      5.56      hexose transport glucose transporter activity*
```

- When the growth rate is restricted to 0.05, what is the average expression level across all genes in the “response to stress” biological process, separately for each limiting nutrient? What about genes in the “protein biosynthesis” biological process? *Hint*: 3 pipes: `filter`, `group_by`, `summarize`.

```
## # A tibble: 6 × 2
##   nutrient meanexp
##   <chr>      <dbl>
## 1 Ammonia    0.943
## 2 Glucose    0.743
## 3 Leucine    0.811
## 4 Phosphate  0.981
## 5 Sulfate    0.743
## 6 Uracil     0.731
```

```
## # A tibble: 6 × 2
##   nutrient meanexp
##   <chr>      <dbl>
## 1 Ammonia   -1.613
## 2 Glucose   -0.691
## 3 Leucine   -0.574
## 4 Phosphate -0.750
## 5 Sulfate   -0.913
## 6 Uracil    -0.880
```

## Exercise set 4

That was easy, right? How about some tougher ones.

First, some review. How do we see the number of distinct values of a variable? Use `n_distinct()` within a `summarize()` call.

```
ydat %>% summarize(n_distinct(mf))
```

```
## # A tibble: 1 × 1
##   `n_distinct(mf)`
##   <int>
## 1         1086
```

- Which 10 biological process annotations have the most genes associated with them? What about molecular functions? *Hint*: 4 pipes: `group_by`, `summarize` with `n_distinct`, `arrange`, `head`.

```
## # A tibble: 10 × 2
##                                     bp      n
##                                     <chr> <int>
## 1 biological process unknown        269
## 2 protein biosynthesis              182
## 3 protein amino acid phosphorylation*  78
## 4 protein biosynthesis*              73
## 5 cell wall organization and biogenesis* 64
## 6 regulation of transcription from RNA polymerase II promoter* 49
## 7 nuclear mRNA splicing, via spliceosome 47
## 8 DNA repair*                        44
## 9 aerobic respiration*                42
## 10 ER to Golgi transport*            42
```

```
## # A tibble: 10 × 2
```

```
##               mf      n
##             <chr> <int>
## 1      molecular function unknown 886
## 2      structural constituent of ribosome 185
## 3              protein binding 107
## 4              RNA binding 63
## 5      protein binding* 53
## 6              DNA binding* 44
## 7      structural molecule activity 43
## 8              GTPase activity 40
## 9      structural constituent of cytoskeleton 39
## 10     transcription factor activity 38
```

2. How many distinct genes are there where we know what process the gene is involved in but we don't know what it does? *Hint*: 3 pipes; filter where `bp!="biological process unknown"` & `mf=="molecular function unknown"`, and after selecting columns of interest, pipe the output to `distinct()`. The answer should be **737**, and here are a few:

```
## # A tibble: 737 × 3
##   symbol                                bp
##   <chr>                                <chr>
## 1   SFB2                                ER to Golgi transport
## 2   EDC3                                deadenylation-independent decapping
## 3   PER1                                response to unfolded protein*
## 4   PEX25                                peroxisome organization and biogenesis*
## 5   BNI5                                cytokinesis*
## 6   CSN12 adaptation to pheromone during conjugation with cellular fusion
## 7   SEC39                                secretory pathway
## 8   ABC1                                ubiquinone biosynthesis
## 9   PRP46                                nuclear mRNA splicing, via spliceosome
## 10  MAM3                                mitochondrion organization and biogenesis*
## # ... with 727 more rows, and 1 more variables: mf <chr>
```

3. When the growth rate is restricted to 0.05 by limiting Glucose, which biological processes are the most upregulated? Show a sorted list with the most upregulated BPs on top, displaying the biological process and the average expression of all genes in that process rounded to two digits. *Hint*: 5 pipes: `filter`, `group_by`, `summarize`, `mutate`, `arrange`.

```
## # A tibble: 881 × 2
##               bp meanexp
##             <chr>   <dbl>
## 1      fermentation* 6.28
## 2      glyoxylate cycle 5.29
## 3 oxygen and reactive oxygen species metabolism 5.04
## 4      fumarate transport* 5.03
## 5      acetyl-CoA biosynthesis* 4.32
## 6      gluconeogenesis 3.64
## 7      fatty acid beta-oxidation 3.57
## 8      lactate transport 3.48
## 9      carnitine metabolism 3.30
## 10     alcohol metabolism* 3.25
## # ... with 871 more rows
```

4. Group the data by limiting nutrient (primarily) then by biological process. Get the average expression for all genes annotated with each process, separately for each limiting nutrient, where the growth rate is restricted to 0.05. Arrange the result to show the most upregulated processes on top. The initial

result will look like the result below. Pipe this output to a `View()` statement. What's going on? Why didn't the `arrange()` work? *Hint*: 5 pipes: `filter`, `group_by`, `summarize`, `arrange`, `View`.

```
## Source: local data frame [5,257 x 3]
## Groups: nutrient [6]
##
##      nutrient                                bp meanexp
##      <chr>                                <chr>    <dbl>
## 1  Ammonia                                allantoate transport    6.64
## 2  Ammonia                                amino acid transport*    6.64
## 3  Phosphate                            glycerophosphodiester transport    6.64
## 4  Glucose                                fermentation*    6.28
## 5  Ammonia                                allantoin transport    5.56
## 6  Glucose                                glyoxylate cycle    5.29
## 7  Ammonia                                proline catabolism*    5.14
## 8  Ammonia                                urea transport    5.14
## 9  Glucose oxygen and reactive oxygen species metabolism    5.04
## 10 Glucose                                fumarate transport*    5.03
## # ... with 5,247 more rows
```

- Let's try to further process that result to get only the top three most upregulated biological processes for each limiting nutrient. Google search "dplyr first result within group." You'll need a `filter(row_number().....)` in there somewhere. *Hint*: 5 pipes: `filter`, `group_by`, `summarize`, `arrange`, `filter(row_number()....` *Note*: dplyr's pipe syntax used to be `%.%` before it changed to `%>%`. So when looking around, you might still see some people use the old syntax. Now if you try to use the old syntax, you'll get a deprecation warning.

```
## Source: local data frame [18 x 3]
## Groups: nutrient [6]
##
##      nutrient                                bp meanexp
##      <chr>                                <chr>    <dbl>
## 1  Ammonia                                allantoate transport    6.64
## 2  Ammonia                                amino acid transport*    6.64
## 3  Phosphate                            glycerophosphodiester transport    6.64
## 4  Glucose                                fermentation*    6.28
## 5  Ammonia                                allantoin transport    5.56
## 6  Glucose                                glyoxylate cycle    5.29
## 7  Glucose oxygen and reactive oxygen species metabolism    5.04
## 8  Uracil                                fumarate transport*    4.32
## 9  Phosphate                            vacuole fusion, non-autophagic    4.20
## 10 Leucine                                fermentation*    4.15
## 11 Phosphate                            regulation of cell redox homeostasis*    4.03
## 12 Leucine                                fumarate transport*    3.72
## 13 Leucine                                glyoxylate cycle    3.65
## 14 Sulfate                                protein ubiquitination    3.40
## 15 Sulfate                                fumarate transport*    3.27
## 16 Uracil                                pyridoxine metabolism    3.11
## 17 Uracil                                asparagine catabolism*    3.06
## 18 Sulfate                                sulfur amino acid metabolism*    2.69
```

- There's a slight problem with the examples above. We're getting the average expression of all the biological processes separately by each nutrient. But some of these biological processes only have a single gene in them! If we tried to do the same thing to get the correlation between rate and expression, the calculation would work, but we'd get a warning about a standard deviation being zero. The correlation coefficient value that results is `NA`, i.e., missing. While we're summarizing the correlation between rate

and expression, let's also show the number of distinct genes within each grouping.

```
ydat %>%
  group_by(nutrient, bp) %>%
  summarize(r=cor(rate, expression), ngenes=n_distinct(symbol))
```

```
## Warning in cor(c(0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05,
## 0.05, : the standard deviation is zero
```

```
## Source: local data frame [5,286 x 4]
## Groups: nutrient [?]
```

```
##
```

	nutrient	bp	r	ngenes
	<chr>	<chr>	<dbl>	<int>
## 1	Ammonia	'de novo' IMP biosynthesis*	0.3125	8
## 2	Ammonia	'de novo' pyrimidine base biosynthesis	-0.0482	3
## 3	Ammonia	'de novo' pyrimidine base biosynthesis*	0.1670	4
## 4	Ammonia	35S primary transcript processing	0.5080	13
## 5	Ammonia	35S primary transcript processing*	0.4240	30
## 6	Ammonia	acetate biosynthesis	0.4677	1
## 7	Ammonia	acetate metabolism	0.9291	1
## 8	Ammonia	acetate metabolism*	-0.6855	1
## 9	Ammonia	acetyl-CoA biosynthesis	-0.8512	1
## 10	Ammonia	acetyl-CoA biosynthesis from pyruvate	0.0951	1

```
## # ... with 5,276 more rows
```

- Take the above code and continue to process the result to show only results where the process has at least 5 genes. Add a column corresponding to the absolute value of the correlation coefficient, and show for each nutrient the singular process with the highest correlation between rate and expression, regardless of direction. *Hint:* 4 more pipes: `filter`, `mutate`, `arrange`, and `filter` again with `row_number()==1`. Ignore the warning.

```
## Source: local data frame [6 x 5]
## Groups: nutrient [6]
```

```
##
```

	nutrient	bp	r	ngenes	absr
	<chr>	<chr>	<dbl>	<int>	<dbl>
## 1	Glucose	telomerase-independent telomere maintenance	-0.95	7	0.95
## 2	Ammonia	telomerase-independent telomere maintenance	-0.91	7	0.91
## 3	Leucine	telomerase-independent telomere maintenance	-0.90	7	0.90
## 4	Phosphate	telomerase-independent telomere maintenance	-0.90	7	0.90
## 5	Uracil	telomerase-independent telomere maintenance	-0.81	7	0.81
## 6	Sulfate	translational elongation*	0.79	5	0.79